

Assignment 8: Time Series Analysis

Sloane Hauck

Fall 2025

OVERVIEW

This exercise accompanies the lessons in Environmental Data Analytics on generalized linear models.

Directions

1. Rename this file `<FirstLast>_A08_TimeSeries.Rmd` (replacing `<FirstLast>` with your first and last name).
2. Change “Student Name” on line 3 (above) with your name.
3. Work through the steps, **creating code and output** that fulfill each instruction.
4. Be sure to **answer the questions** in this assignment document.
5. When you have completed the assignment, **Knit** the text and code into a single PDF file.

Set up

1. Set up your session:
 - Check your working directory
 - Load the tidyverse, lubridate, zoo, and trend packages
 - Set your ggplot theme

```
library(tidyverse);library(zoo);library(lubridate);library(trend);library(here)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    4.0.0      v tibble    3.2.1
## v lubridate  1.9.4      v tidyr     1.3.1
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
##
## Attaching package: 'zoo'
##
##
## The following objects are masked from 'package:base':
##
```

```
##      as.Date, as.Date.numeric
##
##
## here() starts at /home/guest/EDE_Fall2025
```

```
here()
```

```
## [1] "/home/guest/EDE_Fall2025"
```

```
mytheme <- theme_classic(base_size = 14) +
  theme(axis.text = element_text(color = "navy"),
        legend.position = "right",
        panel.background = element_rect(fill = "white")
  )
theme_set(mytheme)
```

2. Import the ten datasets from the Ozone_TimeSeries folder in the Raw data folder. These contain ozone concentrations at Garinger High School in North Carolina from 2010-2019 (the EPA air database only allows downloads for one year at a time). Import these either individually or in bulk and then combine them into a single dataframe named **GaringerOzone** of 3589 observation and 20 variables.

```
#1

files <- list.files(path = "~/EDE_Fall2025/Data/Raw/Ozone_TimeSeries/",
                    pattern = ".csv$",
                    full.names = TRUE)

list_files <- lapply(files, read.csv)

Ozone_Data <- do.call(rbind, list_files)
```

Wrangle

3. Set your date column as a date class.
4. Wrangle your dataset so that it only contains the columns Date, Daily.Max.8.hour.Ozone.Concentration, and DAILY_AQI_VALUE.
5. Notice there are a few days in each year that are missing ozone concentrations. We want to generate a daily dataset, so we will need to fill in any missing days with NA. Create a new data frame that contains a sequence of dates from 2010-01-01 to 2019-12-31 (hint: `as.data.frame(seq())`). Call this new data frame Days. Rename the column name in Days to "Date".
6. Use a `left_join` to combine the data frames. Specify the correct order of data frames within this function so that the final dimensions are 3652 rows and 3 columns. Call your combined data frame GaringerOzone.

```
# 3

class(Ozone_Data$Date)
```

```
## [1] "character"
```

```

Ozone_Data$Date <- mdy(Ozone_Data$Date)

# 4

Ozone_Sub <- Ozone_Data %>%
  select(c(Date, Daily.Max.8.hour.Ozone.Concentration, DAILY_AQI_VALUE))

# 5

Days <- as.data.frame(seq.Date(from = as.Date("2010/01/01"),
                              to = as.Date("2019/12/31"),
                              by = "day"))

colnames(Days) <- "Date"

# 6

GaringerOzone <- Days %>%
  left_join(Ozone_Sub, by = "Date")

```

Visualize

7. Create a line plot depicting ozone concentrations over time. In this case, we will plot actual concentrations in ppm, not AQI values. Format your axes accordingly. Add a smoothed line showing any linear trend of your data. Does your plot suggest a trend in ozone concentration over time?

```


#7

ggplot(GaringerOzone, aes(x = Date, y = Daily.Max.8.hour.Ozone.Concentration)) +
  geom_line() +
  geom_smooth(method = "lm") +
  labs(y = "Daily Max Ozone (ppm)")

## 'geom_smooth()' using formula = 'y ~ x'

## Warning: Removed 63 rows containing non-finite outside the scale range
## ('stat_smooth()').

```



Answer: The slope of the trend line shows a general decrease in ozone concentrations over the entire time period.

Time Series Analysis

Study question: Have ozone concentrations changed over the 2010s at this station?

8. Use a linear interpolation to fill in missing daily data for ozone concentration. Why didn't we use a piecewise constant or spline interpolation?

```
#8

GaringerOzone <- GaringerOzone %>%
  mutate(Daily.Max.8.hour.Ozone.Concentration =
    zoo::na.approx(Daily.Max.8.hour.Ozone.Concentration) )

summary(GaringerOzone$Daily.Max.8.hour.Ozone.Concentration)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.00200 0.03200 0.04100 0.04151 0.05100 0.09300
```

Answer: Because the values are changing daily, we do not want to assume a constant value as the piecewise constant would. The spline interpolation uses a quadratic function to determine the value, but the time periods missing in the Garinger Ozone data are so small, that interpolation would be inappropriate compared to simply “connecting the dots.”

9. Create a new data frame called `GaringerOzone.monthly` that contains aggregated data: mean ozone concentrations for each month. In your pipe, you will need to first add columns for year and month to form the groupings. In a separate line of code, create a new Date column with each month-year combination being set as the first day of the month (this is for graphing purposes only)

#9

```
GaringerOzone_Monthly <- GaringerOzone %>%  
  mutate(Month = month(Date),  
         Year = year(Date),  
         Date = as.yearmon(Date)) %>%  
  group_by(Date) %>%  
  summarize(Monthly.Mean = mean(Daily.Max.8.hour.Ozone.Concentration))
```

10. Generate two time series objects. Name the first `GaringerOzone.daily.ts` and base it on the dataframe of daily observations. Name the second `GaringerOzone.monthly.ts` and base it on the monthly average ozone values. Be sure that each specifies the correct start and end dates and the frequency of the time series.


#10

```
f_day <- day(first(GaringerOzone$Date))  
f_month <- month(first(GaringerOzone$Date))  
f_year <- year(first(GaringerOzone$Date))  
  
GaringerOzone.daily.ts <- ts(GaringerOzone$Daily.Max.8.hour.Ozone.Concentration,  
                             start=c(f_year,f_month, f_day),  
                             frequency=365)  
  
GaringerOzone.monthly.ts <- ts(GaringerOzone_Monthly$Monthly.Mean,  
                               start=c(f_year,f_month),  
                               frequency=12)
```


11. Decompose the daily and the monthly time series objects and plot the components using the `plot()` function.

#11

```
GaringerOzone.daily.decompose <- stl(GaringerOzone.daily.ts,  
                                     s.window = "periodic")  
plot(GaringerOzone.daily.decompose)
```



```
GaringerOzone.monthly.decompose <- stl(GaringerOzone.monthly.ts,
                                       s.window = "periodic")
plot(GaringerOzone.monthly.decompose)
```



12. Run a monotonic trend analysis for the monthly Ozone series. In this case the seasonal Mann-Kendall is most appropriate; why is this?

#12

```
GaringerOzone_Monthly_Trend <- trend::smk.test(GaringerOzone.monthly.ts)
GaringerOzone_Monthly_Trend
```

```
##
## Seasonal Mann-Kendall trend test (Hirsch-Slack test)
##
## data: GaringerOzone.monthly.ts
## z = -1.963, p-value = 0.04965
## alternative hypothesis: true S is not equal to 0
## sample estimates:
## S varS
## -77 1499
```

```
summary(GaringerOzone_Monthly_Trend)
```

```
##
## Seasonal Mann-Kendall trend test (Hirsch-Slack test)
##
## data: GaringerOzone.monthly.ts
```


```
## alternative hypothesis: two.sided
##
## Statistics for individual seasons
##
## H0
##
##      S varS    tau      z Pr(>|z|)
## Season 1:  S = 0   15  125  0.333  1.252  0.21050
## Season 2:  S = 0   -1  125 -0.022  0.000  1.00000
## Season 3:  S = 0   -4  124 -0.090 -0.269  0.78762
## Season 4:  S = 0  -17  125 -0.378 -1.431  0.15241
## Season 5:  S = 0  -15  125 -0.333 -1.252  0.21050
## Season 6:  S = 0  -17  125 -0.378 -1.431  0.15241
## Season 7:  S = 0  -11  125 -0.244 -0.894  0.37109
## Season 8:  S = 0   -7  125 -0.156 -0.537  0.59151
## Season 9:  S = 0   -5  125 -0.111 -0.358  0.72051
## Season 10:  S = 0  -13  125 -0.289 -1.073  0.28313
## Season 11:  S = 0  -13  125 -0.289 -1.073  0.28313
## Season 12:  S = 0   11  125  0.244  0.894  0.37109
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Answer: Ozone fluctuates seasonally, so by applying the seasonal Mann-Kendall test, we ensure that there is no seasonal bias in the data trend.

13. Create a plot depicting mean monthly ozone concentrations over time, with both a `geom_point` and a `geom_line` layer. Edit your axis labels accordingly.

```
# 13

ggplot(GaringerOzone_Monthly, aes(x = as.Date(Date),
                                   y = Monthly.Mean,
                                   group = 1)) +
  geom_line(color = "lightblue") +
  geom_point(color = "navy") +
  scale_x_date(date_labels = "%b %Y",
               breaks = "6 months") +
  labs(y = "Monthly Mean (ppm)",
       x = "Date") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

14. To accompany your graph, summarize your results in context of the research question. Include output from the statistical test in parentheses at the end of your sentence. Feel free to use multiple sentences in your interpretation.

Answer: Based on the monthly data for the seasonal Mann-Kendall test, we see weak evidence of a slightly decreasing seasonal trend. The negative Z-value implies a decreasing trend, and a p-value of 0.04965 is just below 0.05, so it is a weak trend. This means that mean monthly ozone concentrations observed at Garinger are getting decreasing slightly data:

GaringerOzone.monthly.ts z = -1.963, p-value = 0.04965 alternative hypothesis: true S is not equal to 0 sample estimates: S varS -77 1499

15. Subtract the seasonal component from the `GaringerOzone.monthly.ts`. Hint: Look at how we extracted the series components for the `EnoDischarge` on the lesson `Rmd` file.
16. Run the Mann Kendall test on the non-seasonal Ozone monthly series. Compare the results with the ones obtained with the Seasonal Mann Kendall on the complete series.

#15

```
GaringerOzone.Monthly.NoSeason <- GaringerOzone.monthly.ts -
  GaringerOzone.monthly.decompose$time.series[,1]
```

#16

```
GaringerOzone_Monthly_NoSeason_Trend <- trend::mk.test(GaringerOzone.Monthly.NoSeason)

GaringerOzone_Monthly_NoSeason_Trend
```

```
##
## Mann-Kendall trend test
##
## data: GaringerOzone.Monthly.NoSeason
## z = -2.672, n = 120, p-value = 0.00754
## alternative hypothesis: true S is not equal to 0
## sample estimates:
##          S          varS          tau
## -1.179000e+03  1.943657e+05 -1.651376e-01
```

Answer: Similar to the seasonal Mann-Kendall test, we are seeing a decreasing trend as visible in the z-value for the output (-2.672). However, the p-value decreased, showing that the seasonal trend was masking the decreasing trend in the data and by removing it, we can more accurately measure statistical significance in the decreasing monthly mean ozone.